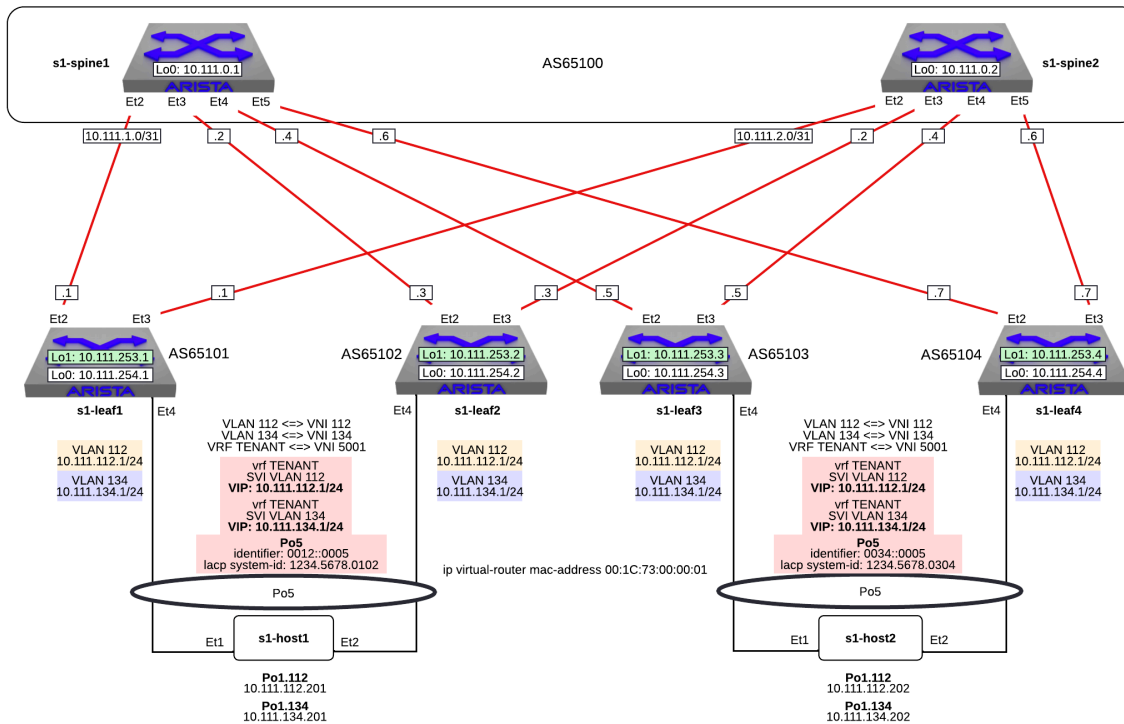


# L2 and L3 EVPN - Symmetric IRB with All-Active Multihoming



(\_images/nested\_l2l3evpn-aa\_topo\_dual\_dc.png)

## Note

This lab exercise is focused on the VXLAN EVPN configuration. IP addresses and BGP Underlay are already configured.

- Log into the **LabAccess** jumpserver:
  - Type `97` to access additional lab, then `evpn-labs` at the prompt to access the EVPN VXLAN content. Then type `l2l3evpn-aa` for the Layer 2 and 3 EVPN lab. The script will configure the datacenter with the exception of **s1-leaf4**.

**Note**

Did you know the “l2l3evpn-aa” script is composed of Python code that uses the CloudVision Portal REST API to automate the provisioning of CVP Configlets. The configlets that are configured via the REST API are

L2L3EVPN-AA\_s1-spine1 , L2L3EVPN-AA\_s1-spine2 , L2L3EVPN-AA\_s1-leaf1 ,  
L2L3EVPN-AA\_s1-leaf2 , L2L3EVPN-AA\_s1-leaf3 , L2L3EVPN-AA\_s1-leaf4 .

2. On **s1-leaf4**, check if Multi-Agent Routing Protocols are enabled.

```
s1-leaf4#show run section service
service routing protocols model multi-agent
s1-leaf4#show ip route summary
```

```
Operating routing protocol model: multi-agent
Configured routing protocol model: multi-agent
```

VRF: default

Route Source	Number Of Routes
connected	4
static (persistent)	0
static (non-persistent)	0
VXLAN Control Service	0
static nexthop-group	0
ospf	0
Intra-area: 0 Inter-area: 0 External-1: 0 External-2: 0	
NSSA External-1: 0 NSSA External-2: 0	
ospfv3	0
bgp	9
External: 7 Internal: 2	
isis	0
Level-1: 0 Level-2: 0	
rip	0
internal	11
attached	3
aggregate	0
dynamic policy	0
gribi	0
Total Routes	27

Number of routes per mask-length:

/8: 2	/24: 3	/30: 1	/31: 2	/32: 19
-------	--------	--------	--------	---------

**Note**

By default, EOS is using the GateD routing process. Activating (ArBGP) is requiring a reboot. This has been done prior to the lab buildout so no reboot is required here.

3. On **s1-leaf4**, check the following operational states before configuring EVPN constructs:

a. Verify BGP operational details for Underlay:

#### Note

You should see underlay sessions; one to each spine. In this design, there is no “peer-link”

```
s1-leaf4#show ip bgp summary
BGP summary information for VRF default
Router identifier 10.111.254.4, local AS number 65102
Neighbor Status Codes: m - Under maintenance
```

Neighbor	V AS	MsgRcvd	MsgSent	InQ	OutQ	Up/Down	State	PfxRcd	PfxAcc
10.111.1.6	4 65100	9	12	0	0	00:00:07	Estab	8	8
10.111.2.6	4 65100	9	12	0	0	00:00:07	Estab	8	8

b. Check the IP routing table:

#### Note

Notice that **s1-leaf4** has 2 ECMP paths for reaching **s1-leaf1**, **s1-leaf2** and **s1-leaf3** loopbacks.

```
s1-leaf4#show ip route
```

```
VRF: default
```

```
Codes: C - connected, S - static, K - kernel,
```

```
    O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
```

```
    E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
```

```
    N2 - OSPF NSSA external type2, B - Other BGP Routes,
```

```
    B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
```

```
    I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
```

```
    A O - OSPF Summary, NG - Nexthop Group Static Route,
```

```
    V - VXLAN Control Service, M - Martian,
```

```
    DH - DHCP client installed default route,
```

```
    DP - Dynamic Policy Route, L - VRF Leaked,
```

```
    G - gRIBI, RC - Route Cache Route
```

```
Gateway of last resort is not set
```

```
B E      10.111.0.1/32 [200/0] via 10.111.1.6, Ethernet2
```

```
B E      10.111.0.2/32 [200/0] via 10.111.2.6, Ethernet3
```

```
C        10.111.1.6/31 is directly connected, Ethernet2
```

```
B E      10.111.1.0/24 [200/0] via 10.111.1.6, Ethernet2
```

```
C        10.111.2.6/31 is directly connected, Ethernet3
```

```
B E      10.111.2.0/24 [200/0] via 10.111.2.6, Ethernet3
```

```
B I      10.111.112.0/24 [200/0] via 10.255.255.1, Vlan4094
```

```
B E      10.111.253.1/32 [200/0] via 10.111.1.6, Ethernet2
                                     via 10.111.2.6, Ethernet3
```

```
B E      10.111.253.2/32 [200/0] via 10.111.1.6, Ethernet2
                                     via 10.111.2.6, Ethernet3
```

```
B E      10.111.253.3/32 [200/0] via 10.111.1.6, Ethernet2
                                     via 10.111.2.6, Ethernet3
```

```
B E      10.111.254.1/32 [200/0] via 10.111.1.6, Ethernet2
                                     via 10.111.2.6, Ethernet3
```

```
B E      10.111.254.2/32 [200/0] via 10.111.1.6, Ethernet2
                                     via 10.111.2.6, Ethernet3
```

```
B E      10.111.254.3/32 [200/0] via 10.111.1.6, Ethernet2
                                     via 10.111.2.6, Ethernet3
```

```
C        10.111.254.4/32 is directly connected, Loopback0
```

```
C        10.255.255.0/30 is directly connected, Vlan4094
```

```
C        192.168.0.0/24 is directly connected, Management0
```

#### 4. On **s1-leaf4**, configure the BGP EVPN control-plane.

##### a. Configure the EVPN control plane.

#### Note

In this lab, the Spines serve as EVPN Route Servers. They receive the EVPN Routes from each leaf and, due to our eBGP setup, will naturally pass them along the other leaves. In an EVPN A-A setup with eBGP, each VTEP has its own unique ASN.

Also note that BGP standard and extended communities are explicitly enabled on the peering. EVPN makes use of extended BGP communities for route signaling and standard communities allow for various other functions such as BGP maintenance mode.

Lastly, note in this setup we use eBGP-multihop peerings with the Loopback0 interfaces of each switch. This follows Arista best-practice designs for separation of Underlay (peerings done using physical Ethernet interfaces) and Overlay (peerings done using Loopbacks) when leveraging eBGP. Other options exist and can be discussed with your Arista SE.

```
router bgp 65104
  neighbor SPINE-EVPN peer group
  neighbor SPINE-EVPN remote-as 65100
  neighbor SPINE-EVPN update-source Loopback0
  neighbor SPINE-EVPN ebgp-multihop 3
  neighbor SPINE-EVPN send-community standard extended
  neighbor 10.111.0.1 peer group SPINE-EVPN
  neighbor 10.111.0.2 peer group SPINE-EVPN
  !
  address-family evpn
    neighbor SPINE-EVPN activate
```

- b. Verify the EVPN Control-Plane is established to both Spine peers.

```
s1-leaf4(config-router-bgp-af)#show bgp evpn summary
BGP summary information for VRF default
Router identifier 10.111.254.4, local AS number 65104
Neighbor Status Codes: m - Under maintenance
```

Neighbor	V	AS	MsgRcvd	MsgSent	InQ	OutQ	Up/Down	State	PfxRcd	PfxAcc
10.111.0.1	4	65100	31	31	0	0	00:00:04	Estab	34	34
10.111.0.2	4	65100	31	4	0	0	00:00:04	Estab	34	34

5. On **s1-leaf4**, configure the VXLAN data-plane for transport.

- a. Configure Loopback1 with the shared IP of **s1-leaf3**.

#### Note

Unlike with MLAG VTEPs, with EVPN A-A, all VTEPs have a unique IP. We will see later how resiliency and load-balancing differ in this setup.

```
interface Loopback1
  description VTEP
  ip address 10.111.253.4/32
```

- b. Configure the Vxlan1 interface with the Loopback1 as the source.

#### Note

This is the logical interface that will provide VXLAN header encap and decap functions.

```
interface Vxlan1
  vxlan source-interface Loopback1
```

6. Configure Layer 2 EVPN services on **s1-leaf4**.

- a. Add the local Layer 2 VLANs with an IDs of 112 and 134.

```
vlan 112
  name Host_Network_112
!
vlan 134
  name Host_Network_134
```

- b. Map the local Layer 2 VLANs with a matching VNIs.

#### Note

This is how the switch understands which local Layer 2 VLAN maps to which VNI in the overlay. The example shows matching them one to one, but any scheme or method is valid, such as adding 10000 to the VLAN ID.

```
interface Vxlan1
  vxlan vlan 112 vni 112
  vxlan vlan 134 vni 134
```

- c. Add the mac-vrf EVPN configuration for VLAN 112 and 134.

#### Note

Here we configure a VLAN-based service with EVPN. It has two components. The first is a route-distinguisher, or **RD** to identify the router (or leaf switch) that is originating the EVPN routes. This can be manually defined in the format of **Number : Number**, such as **Loopback0 : VLAN ID** or as we do in this case, let EOS automatically allocate one.

Second is the route-target, or **RT**. The **RT** is used by the leaf switches in the network to determine if they should import the advertised route into their local table(s). If they receive an EVPN route, they check the **RT** value and see if they have a matching **RT** configured in BGP. If they do, they import the route into the associated mac-vrf (or VLAN). If they do not, they ignore the route.

```
router bgp 65104
!
  vlan 112
    rd auto
    route-target both 112:112
    redistribute learned
!
  vlan 134
    rd auto
    route-target both 134:134
    redistribute learned
```

## 7. Configure Layer 3 EVPN services on **s1-leaf4**.

- a. Create the VRF, or logical routing instance, for the Tenant Layer 3 Network.

#### Note

In EOS, by default, VRFs are created with inter-subnet routing disabled. Always be sure to enable IP routing in user-defined VRFs.

```
vrf instance TENANT
!  
ip routing vrf TENANT
```

- b. Create the SVI for default gateway function for the host network as an Anycast Gateway.

#### Note

With VXLAN, we can leverage a shared IP using Anycast Gateway. This allows a single IP to be shared without any other dedicated IPs per switch.

```
ip virtual-router mac-address 00:1C:73:00:00:01
!  
interface Vlan112  
  description Host Network 112  
  vrf TENANT  
  ip address virtual 10.111.112.1/24  
!  
interface Vlan134  
  description Host Network 134  
  vrf TENANT  
  ip address virtual 10.111.134.1/24
```

- c. Map the local Layer 3 VRF with a matching VNI.

#### Note

For the Layer 3 Service, the VRF requires what is referred to as the Layer 3 VNI, which is used for VXLAN Routing in a Symmetric IRB deployment between VTEPs. Any unique ID number will serve here.

```
interface Vxlan1  
  vxlan vrf TENANT vni 5001
```

- d. Add the IP VRF EVPN configuration for the TENANT VRF.

#### Note

Here we configure a Layer 3 VRF service with EVPN. It also leverage a unique **RD** and **RT**. They are used by the leaf switches for the same purpose as the Layer 2 service. The difference is simply the routes are imported. If they receive a Type 5 EVPN route, they check the **RT** value and see if they have a matching **RT** configured for the VRF. If so, they import the route into the associated VRF routing table. If they do not, they ignore the route.

```

router bgp 65104
  rd auto
  !
  vrf TENANT
    route-target import evpn 5001:5001
    route-target export evpn 5001:5001
    redistribute connected

```

- e. Configure the host-facing EVPN A-A Port-Channel.

### Note

This is where we configure the Ethernet Segment Identifier, or **ESI**, as well as a **RT** value for the Ethernet Segment. We will see how the EVPN control-plane leverages these to negotiate the characteristics and state of the A-A Port-Channel. We also configure a static LACP System-ID. This is to ensure that all members of the Ethernet Segment appear as one LACP system to the downstream device. Note that all these values must match on members of the same Ethernet Segment (or Port-Channel).

```

interface Port-Channel5
  description EVPN A-A Downlink - s1-host2
  switchport trunk allowed vlan 112,134
  switchport mode trunk
  !
  evpn ethernet-segment
    identifier 0034:0000:0000:0000:0005
    route-target import 00:03:04:00:00:05
    lacp system-id 1234.5678.0304
  !
interface Ethernet4
  description EVPN A-A Downlink - s1-host2
  channel-group 5 mode active

```

8. With the Layer 2 and 3 EVPN Services configured, verify the operational state.

- a. Check the VXLAN data-plane configuration on **s1-leaf4**.

### Note

Here we can see some useful commands for VXLAN verification. `show vxlan config-sanity detail` verifies a number of standard things locally and with the MLAG peer to ensure all basic criteria are met. `show interfaces Vxlan1` provides a consolidated series of outputs of operational VXLAN data such as control-plane mode (EVPN in this case), VLAN to VNI mappings and discovered VTEPs.



```
s1-leaf4#show vxlan config-sanity detail
```

Category	Result	Detail
Local VTEP Configuration Check	OK	
Loopback IP Address	OK	
VLAN-VNI Map	OK	
Routing	OK	
VNI VRF ACL	OK	
Decap VRF-VNI Map	OK	
VRF-VNI Dynamic VLAN	OK	
Remote VTEP Configuration Check	OK	
Remote VTEP	OK	
Platform Dependent Check	OK	
VXLAN Bridging	OK	
VXLAN Routing	OK	
CVX Configuration Check	OK	
CVX Server	OK	Not in controller client mode
MLAG Configuration Check	OK	Run 'show mlag config-sanity' to verify MLAG
Peer VTEP IP	OK	MLAG peer is not connected
MLAG VTEP IP	OK	
Peer VLAN-VNI	OK	
Virtual VTEP IP	OK	
MLAG Inactive State	OK	

```
s1-leaf4#show interfaces Vxlan1
Vxlan1 is up, line protocol is up (connected)
Hardware is Vxlan
Source interface is Loopback1 and is active with 10.111.253.4
Replication/Flood Mode is headend with Flood List Source: EVPN
Remote MAC learning via EVPN
VNI mapping to VLANs
Static VLAN to VNI mapping is
  [112, 112]      [134, 134]
Dynamic VLAN to VNI mapping for 'evpn' is
  [4094, 5001]
Note: All Dynamic VLANs used by VCS are internal VLANs.
      Use 'show vxlan vni' for details.
Static VRF to VNI mapping is
  [TENANT, 5001]
Headend replication flood vtep list is:
  112 10.111.253.1  10.111.253.3  10.111.253.2
  134 10.111.253.1  10.111.253.3  10.111.253.2
Shared Router MAC is 0000.0000.0000
```

- b. Determine who the Designated Forwarder is for the EVPN A-A Port-Channel on **s1-leaf4**.

### Note

In an EVPN A-A Ethernet Segment, only one member of the **ES** is elected as the Designated Forwarder, or **DF**. The **DF** is responsible for forwarding BUM traffic to the connected downstream device. By default, a modulus operation is run by all members of the **ES** to uniformly

elect the DF based on the received **Ethernet Segment**, or EVPN Type 4, routes. Highlighted below we can see the received EVPN Type 4 routes from **s1-leaf3** with the matching **ESI** value. The detailed output shows the associated **ES RT** value as well.

By further inspecting the EVPN Instances, or MAC-VRFs, we can determine which member of the **ES** has been elected as the **DF**.

```

s1-leaf4#show bgp evpn route-type ethernet-segment
BGP routing table information for VRF default
Router identifier 10.111.254.4, local AS number 65104
Route status codes: * - valid, > - active, S - Stale, E - ECMP head, e - ECMP
                    c - Contributing to ECMP, % - Pending BGP convergence
Origin codes: i - IGP, e - EGP, ? - incomplete
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link Local

      Network                Next Hop                Metric LocPref Weight Path
* >Ec  RD: 10.111.253.1:1 ethernet-segment 0012:0000:0000:0000:0005 10.111.253.1
      10.111.253.1                -                100      0      65100 65
*   ec  RD: 10.111.253.1:1 ethernet-segment 0012:0000:0000:0000:0005 10.111.253.1
      10.111.253.1                -                100      0      65100 65
* >Ec  RD: 10.111.253.2:1 ethernet-segment 0012:0000:0000:0000:0005 10.111.253.2
      10.111.253.2                -                100      0      65100 65
*   ec  RD: 10.111.253.2:1 ethernet-segment 0012:0000:0000:0000:0005 10.111.253.2
      10.111.253.2                -                100      0      65100 65
* >Ec  RD: 10.111.253.3:1 ethernet-segment 0034:0000:0000:0000:0005 10.111.253.3
      10.111.253.3                -                100      0      65100 65
*   ec  RD: 10.111.253.3:1 ethernet-segment 0034:0000:0000:0000:0005 10.111.253.3
      10.111.253.3                -                100      0      65100 65
* >   RD: 10.111.253.4:1 ethernet-segment 0034:0000:0000:0000:0005 10.111.253.4
      -                            -                -        0      i

s1-leaf4#show bgp evpn route-type ethernet-segment esi 0034:0000:0000:0000:0005 detail
BGP routing table information for VRF default
Router identifier 10.111.254.4, local AS number 65104
BGP routing table entry for ethernet-segment 0034:0000:0000:0000:0005 10.111.253.3, Routes
Paths: 2 available
  65100 65103
    10.111.253.3 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: TunnelEncap:tunnelTypeVxlan EvpnEsImportRt:00:03:04:00:00:05
  65100 65103
    10.111.253.3 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: TunnelEncap:tunnelTypeVxlan EvpnEsImportRt:00:03:04:00:00:05
BGP routing table entry for ethernet-segment 0034:0000:0000:0000:0005 10.111.253.4, Routes
Paths: 1 available
  Local
    - from - (0.0.0.0)
      Origin IGP, metric -, localpref -, weight 0, valid, local, best
      Extended Community: TunnelEncap:tunnelTypeVxlan EvpnEsImportRt:00:03:04:00:00:05
s1-leaf4#show bgp evpn instance
EVPN instance: VLAN 112
Route distinguisher: 0:0
Route target import: Route-Target-AS:112:112
Route target export: Route-Target-AS:112:112
Service interface: VLAN-based
Local VXLAN IP address: 10.111.253.4
VXLAN: enabled
MPLS: disabled
Local ethernet segment:

```

```

ESI: 0034:0000:0000:0000:0005
Interface: Port-Channel5
Mode: all-active
State: up
ES-Import RT: 00:03:04:00:00:05
DF election algorithm: modulus
Designated forwarder: 10.111.253.3
Non-Designated forwarder: 10.111.253.4
EVPN instance: VLAN 134
Route distinguisher: 0:0
Route target import: Route-Target-AS:134:134
Route target export: Route-Target-AS:134:134
Service interface: VLAN-based
Local VXLAN IP address: 10.111.253.4
VXLAN: enabled
MPLS: disabled
Local ethernet segment:
ESI: 0034:0000:0000:0000:0005
Interface: Port-Channel5
Mode: all-active
State: up
ES-Import RT: 00:03:04:00:00:05
DF election algorithm: modulus
Designated forwarder: 10.111.253.3
Non-Designated forwarder: 10.111.253.4

```

- c. On **s1-leaf1**, verify the IMET table to ensure **s1-leaf4** has been discovered in the overlay.

### Note

The Inclusive Multicast Ethernet Tag, or **IMET**, route is how a VTEP advertises membership in a given Layer 2 service, or VXLAN segment. This is also known as the EVPN Type 3 Route. Other leaves receive this route, evaluate the **RT** to see if they have a matching configuration and, if so, import the advertising VTEP into their flood list for BUM traffic. Note that these are done on a per VLAN basis based on the MAC-VRF configuration. Highlighted below are the EVPN Type 3 Routes from **s1-leaf4** which we identify based on the **RD** value. The detail outputs show **RT** and **VNI** information as well as the **Tunnel ID** which in our case is the VTEP address to flood BUM traffic to.

```
s1-leaf1#show bgp evpn route-type imet
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
Route status codes: * - valid, > - active, S - Stale, E - ECMP head, e - ECMP
                   c - Contributing to ECMP, % - Pending BGP convergence
Origin codes: i - IGP, e - EGP, ? - incomplete
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link Local
```

	Network	Next Hop	Metric	LocPref	Weight	Path
* >Ec	RD: 10.111.254.2:112	imet 10.111.253.2				
		10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.254.2:112	imet 10.111.253.2				
		10.111.253.2	-	100	0	65100 65
* >Ec	RD: 10.111.254.2:134	imet 10.111.253.2				
		10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.254.2:134	imet 10.111.253.2				
		10.111.253.2	-	100	0	65100 65
* >Ec	RD: 10.111.254.3:112	imet 10.111.253.3				
		10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.254.3:112	imet 10.111.253.3				
		10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.254.3:134	imet 10.111.253.3				
		10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.254.3:134	imet 10.111.253.3				
		10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.254.4:112	imet 10.111.253.4				
		10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.254.4:112	imet 10.111.253.4				
		10.111.253.4	-	100	0	65100 65
* >Ec	RD: 10.111.254.4:134	imet 10.111.253.4				
		10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.254.4:134	imet 10.111.253.4				
		10.111.253.4	-	100	0	65100 65
* >	RD: 10.111.254.1:112	imet 10.111.253.1				
		-	-	-	0	i
* >	RD: 10.111.254.1:134	imet 10.111.253.1				
		-	-	-	0	i

```
s1-leaf1#show bgp evpn route-type imet rd 10.111.254.4:112 detail
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
BGP routing table entry for imet 10.111.253.4, Route Distinguisher: 10.111.254.4:112
Paths: 2 available
65100 65104
  10.111.253.4 from 10.111.0.2 (10.111.0.2)
    Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
    Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
    VNI: 112
    PMSI Tunnel: Ingress Replication, MPLS Label: 112, Leaf Information Required: false
65100 65104
  10.111.253.4 from 10.111.0.1 (10.111.0.1)
    Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
    Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
```

```

VNI: 112
PMSI Tunnel: Ingress Replication, MPLS Label: 112, Leaf Information Required: fal
s1-leaf1#show interfaces Vxlan1
Vxlan1 is up, line protocol is up (connected)
  Hardware is Vxlan
  Source interface is Loopback1 and is active with 10.111.253.1
  Replication/Flood Mode is headend with Flood List Source: EVPN
  Remote MAC learning via EVPN
  VNI mapping to VLANs
  Static VLAN to VNI mapping is
    [112, 112]          [134, 134]
  Dynamic VLAN to VNI mapping for 'evpn' is
    [4093, 5001]
  Note: All Dynamic VLANs used by VCS are internal VLANs.
        Use 'show vxlan vni' for details.
  Static VRF to VNI mapping is
    [TENANT, 5001]
  Headend replication flood vtep list is:
    112 10.111.253.3    10.111.253.4    10.111.253.2
    134 10.111.253.3    10.111.253.4    10.111.253.2
  Shared Router MAC is 0000.0000.0000

```

- d. Log into **s1-host1** and ping **s2-host2** in both VLANs to populate the network's MAC and ARP tables.

#### Note

Since we are hosting multiple networks on the simulated Hosts, we have separated the networks by VRFs. These are not related to the VRFs in the network fabric. Note that due to host discovery and control-plan convergence in our simulated EOS labs, you may receive some duplicate responses in the initial run. This is normal and should level off upon subsequent ping tests.

```

s1-host1#ping vrf 112 10.111.112.202
PING 10.111.112.202 (10.111.112.202) 72(100) bytes of data.
 80 bytes from 10.111.112.202: icmp_seq=1 ttl=64 time=21.3 ms
 80 bytes from 10.111.112.202: icmp_seq=2 ttl=64 time=17.6 ms
 80 bytes from 10.111.112.202: icmp_seq=3 ttl=64 time=22.2 ms
 80 bytes from 10.111.112.202: icmp_seq=4 ttl=64 time=22.3 ms
 80 bytes from 10.111.112.202: icmp_seq=5 ttl=64 time=23.8 ms

--- 10.111.112.202 ping statistics ---
 5 packets transmitted, 5 received, 0% packet loss, time 64ms
 rtt min/avg/max/mdev = 17.698/21.491/23.822/2.059 ms, pipe 3, ipg/ewma 16.095/21.549 ms
s1-host1#ping vrf 134 10.111.134.202
PING 10.111.134.202 (10.111.134.202) 72(100) bytes of data.
 80 bytes from 10.111.134.202: icmp_seq=1 ttl=64 time=138 ms
 80 bytes from 10.111.134.202: icmp_seq=2 ttl=64 time=132 ms
 80 bytes from 10.111.134.202: icmp_seq=3 ttl=64 time=124 ms
 80 bytes from 10.111.134.202: icmp_seq=4 ttl=64 time=111 ms
 80 bytes from 10.111.134.202: icmp_seq=5 ttl=64 time=103 ms

--- 10.111.134.202 ping statistics ---
 5 packets transmitted, 5 received, 0% packet loss, time 46ms
 rtt min/avg/max/mdev = 103.152/122.104/138.805/13.201 ms, pipe 5, ipg/ewma 11.627/129.4

```

- e. On **s1-leaf1**, check the EVPN control-plane for the associated host MAC/IP.

#### Note

We see the MAC of **s1-host2** multiple times in the control-plane due to our redundant MLAG and ECMP design. Both **s1-leaf3** and **s1-leaf4** are attached to **s1-host2** in VLANs 112 and 134 and therefore will generate these Type 2 EVPN route for its MAC once the host is discovered. They each then send this route up to the redundant Spines (or EVPN Route Servers) which provides an ECMP path to the host. The highlighting below is focusing on **s1-leaf4**. Depending on how traffic hashes from the host, notice that you might **not** see certain entries generated from **s1-leaf4**. This is expected and we will see how aliasing allows the network to understand that the EVPN A-A provides connectivity to the host from each leaf in the ES, whether or not they've individually advertised the host MAC.

Also notice that since we have configured our network for VXLAN Routing functionality we also see the host MAC-IP route that advertises the ARP binding of **s1-host2**. By looking at the detailed output of the command specifically for the host in VNI (VLAN) 112, we can see details about the **RT** and **VNIs**, both Layer 2 (112) and Layer 3 (5001) which we see in further outputs later.

Also highlighted is the ESI value in each Type 2 Route. This signals to the VTEPs that the MAC was learned as part of an EVPN A-A link.

```
s1-leaf1#show bgp evpn route-type mac-ip
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
Route status codes: * - valid, > - active, S - Stale, E - ECMP head, e - ECMP
                   c - Contributing to ECMP, % - Pending BGP convergence
Origin codes: i - IGP, e - EGP, ? - incomplete
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link Local
```

	Network	Next Hop	Metric	LocPref	Weight	Path
* >	RD: 10.111.254.1:112	mac-ip 001c.73c0.c616	-	-	0	i
* >	RD: 10.111.254.1:134	mac-ip 001c.73c0.c616	-	-	0	i
* >	RD: 10.111.254.1:112	mac-ip 001c.73c0.c616 10.111.112.201	-	-	0	i
* >Ec	RD: 10.111.254.2:112	mac-ip 001c.73c0.c616 10.111.112.201	10.111.253.2	-	100	0 65100 65
* ec	RD: 10.111.254.2:112	mac-ip 001c.73c0.c616 10.111.112.201	10.111.253.2	-	100	0 65100 65
* >	RD: 10.111.254.1:134	mac-ip 001c.73c0.c616 10.111.134.201	-	-	0	i
* >Ec	RD: 10.111.254.2:134	mac-ip 001c.73c0.c616 10.111.134.201	10.111.253.2	-	100	0 65100 65
* ec	RD: 10.111.254.2:134	mac-ip 001c.73c0.c616 10.111.134.201	10.111.253.2	-	100	0 65100 65
* >Ec	RD: 10.111.254.3:112	mac-ip 001c.73c0.c617	10.111.253.3	-	100	0 65100 65
* ec	RD: 10.111.254.3:112	mac-ip 001c.73c0.c617	10.111.253.3	-	100	0 65100 65
* >Ec	RD: 10.111.254.3:134	mac-ip 001c.73c0.c617	10.111.253.3	-	100	0 65100 65
* ec	RD: 10.111.254.3:134	mac-ip 001c.73c0.c617	10.111.253.3	-	100	0 65100 65
* >Ec	RD: 10.111.254.3:112	mac-ip 001c.73c0.c617 10.111.112.202	10.111.253.3	-	100	0 65100 65
* ec	RD: 10.111.254.3:112	mac-ip 001c.73c0.c617 10.111.112.202	10.111.253.3	-	100	0 65100 65
* >Ec	RD: 10.111.254.4:112	mac-ip 001c.73c0.c617 10.111.112.202	10.111.253.4	-	100	0 65100 65
* ec	RD: 10.111.254.4:112	mac-ip 001c.73c0.c617 10.111.112.202	10.111.253.4	-	100	0 65100 65
* >Ec	RD: 10.111.254.3:134	mac-ip 001c.73c0.c617 10.111.134.202	10.111.253.3	-	100	0 65100 65
* ec	RD: 10.111.254.3:134	mac-ip 001c.73c0.c617 10.111.134.202	10.111.253.3	-	100	0 65100 65
* >Ec	RD: 10.111.254.4:134	mac-ip 001c.73c0.c617 10.111.134.202	10.111.253.4	-	100	0 65100 65
* ec	RD: 10.111.254.4:134	mac-ip 001c.73c0.c617 10.111.134.202	10.111.253.4	-	100	0 65100 65

```
s1-leaf1#show bgp evpn route-type mac-ip 001c.73c0.c617 vni 112 detail
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
```



```

BGP routing table entry for mac-ip 001c.73c0.c617, Route Distinguisher: 10.111.254.3:11
Paths: 2 available
 65100 65103
    10.111.253.3 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
      VNI: 112 ESI: 0034:0000:0000:0000:0005
 65100 65103
    10.111.253.3 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
      VNI: 112 ESI: 0034:0000:0000:0000:0005
BGP routing table entry for mac-ip 001c.73c0.c617 10.111.112.202, Route Distinguisher:
Paths: 2 available
 65100 65103
    10.111.253.3 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: Route-Target-AS:112:112 Route-Target-AS:5001:5001 TunnelEncap
      VNI: 112 L3 VNI: 5001 ESI: 0034:0000:0000:0000:0005
 65100 65103
    10.111.253.3 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: Route-Target-AS:112:112 Route-Target-AS:5001:5001 TunnelEncap
      VNI: 112 L3 VNI: 5001 ESI: 0034:0000:0000:0000:0005
BGP routing table entry for mac-ip 001c.73c0.c617 10.111.112.202, Route Distinguisher:
Paths: 2 available
 65100 65104
    10.111.253.4 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: Route-Target-AS:112:112 Route-Target-AS:5001:5001 TunnelEncap
      VNI: 112 L3 VNI: 5001 ESI: 0034:0000:0000:0000:0005
 65100 65104
    10.111.253.4 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: Route-Target-AS:112:112 Route-Target-AS:5001:5001 TunnelEncap
      VNI: 112 L3 VNI: 5001 ESI: 0034:0000:0000:0000:0005

```

1. On **s1-leaf1**, check the EVPN control-plane for the EVPN A-A Signaling associated with the **s1-host2**.

### Note

We saw above that the Type 2 routes contained an **ESI** value. We can then determine all of the VTEPs that are members of that **ES** by inspecting the **Auto-Discovery**, or EVPN Type 1, routes. Highlighted below are the entries associated with the EVPN A-A **ES** that is attached to **s1-host2**. **s1-leaf1** has learned that both **s1-leaf3** and **s1-leaf4** are members of the same **ES**. This is done on a per MAC-VRF (or VLAN) basis.

By looking at the detailed output for that **ESI** specifically for VNI 112, we can see further information about associated **RT** and **VNI** information. By interpreting this, **s1-leaf1** understands that to reach **s1-host2**, packets can be sent to either **s1-leaf3** OR **s1-leaf4** since they are mem-

bers of the same **ES** where the **s1-host2** is attached (even though **s1-lea4** never generated a Type 2 MAC Only route in our example).

```
s1-leaf1#show bgp evpn route-type auto-discovery
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
Route status codes: * - valid, > - active, S - Stale, E - ECMP head, e - ECMP
                   c - Contributing to ECMP, % - Pending BGP convergence
Origin codes: i - IGP, e - EGP, ? - incomplete
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link Local
```

	Network	Next Hop	Metric	LocPref	Weight	Path
* >	RD: 10.111.254.1:112	auto-discovery 0	0012:0000:0000:0000:0005			
		-	-	-	0	i
* >	RD: 10.111.254.1:134	auto-discovery 0	0012:0000:0000:0000:0005			
		-	-	-	0	i
* >Ec	RD: 10.111.254.2:112	auto-discovery 0	0012:0000:0000:0000:0005			
		10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.254.2:112	auto-discovery 0	0012:0000:0000:0000:0005			
		10.111.253.2	-	100	0	65100 65
* >Ec	RD: 10.111.254.2:134	auto-discovery 0	0012:0000:0000:0000:0005			
		10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.254.2:134	auto-discovery 0	0012:0000:0000:0000:0005			
		10.111.253.2	-	100	0	65100 65
* >	RD: 10.111.253.1:1	auto-discovery 0	0012:0000:0000:0000:0005			
		-	-	-	0	i
* >Ec	RD: 10.111.253.2:1	auto-discovery 0	0012:0000:0000:0000:0005			
		10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.253.2:1	auto-discovery 0	0012:0000:0000:0000:0005			
		10.111.253.2	-	100	0	65100 65
* >Ec	RD: 10.111.254.3:112	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.254.3:112	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.254.3:134	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.254.3:134	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.254.4:112	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.254.4:112	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.4	-	100	0	65100 65
* >Ec	RD: 10.111.254.4:134	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.254.4:134	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.4	-	100	0	65100 65
* >Ec	RD: 10.111.253.3:1	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.253.3:1	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.253.4:1	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.253.4:1	auto-discovery 0	0034:0000:0000:0000:0005			
		10.111.253.4	-	100	0	65100 65

```
s1-leaf1#show bgp evpn route-type auto-discovery vni 112 esi 0034:0000:0000:0000:0005 d
```

```

BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
BGP routing table entry for auto-discovery 0 0034:0000:0000:0000:0005, Route Distinguis
Paths: 2 available
 65100 65103
    10.111.253.3 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
      VNI: 112
 65100 65103
    10.111.253.3 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
      VNI: 112
BGP routing table entry for auto-discovery 0 0034:0000:0000:0000:0005, Route Distinguis
Paths: 2 available
 65100 65104
    10.111.253.4 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
      VNI: 112
 65100 65104
    10.111.253.4 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: Route-Target-AS:112:112 TunnelEncap:tunnelTypeVxlan
      VNI: 112

```

- f. On **s1-leaf1**, verify the BGP table to ensure the Tenant networks on **s1-leaf4** has been learned in the overlay.

### Note

The output below shows learned **IP Prefix** routes from EVPN. These are referred to as EVPN Type 5 routes. Similar to the Type 2 and 3 Routes, other VTEPs evaluate the **RT** to see if they have a matching configuration and, if so, import the contained prefix into their VRF Route Table. Note that IPv4 and IPv6 are supported.

In the detailed output, we can see the specific routes from **s1-leaf4** by filtering based on the **RD** value. We can see information about the **RT**, EVPN Router MAC (shared with **s1-leaf3**) and the L3 VNI. The highlights below focus on the 10.111.112.0/24 network.

```
s1-leaf1#show bgp evpn route-type ip-prefix ipv4
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
Route status codes: * - valid, > - active, S - Stale, E - ECMP head, e - ECMP
                   c - Contributing to ECMP, % - Pending BGP convergence
Origin codes: i - IGP, e - EGP, ? - incomplete
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link Local
```

	Network	Next Hop	Metric	LocPref	Weight	Path
* >	RD: 10.111.254.1:1 ip-prefix 10.111.112.0/24	-	-	-	0	i
* >Ec	RD: 10.111.254.2:1 ip-prefix 10.111.112.0/24	10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.254.2:1 ip-prefix 10.111.112.0/24	10.111.253.2	-	100	0	65100 65
* >Ec	RD: 10.111.254.3:1 ip-prefix 10.111.112.0/24	10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.254.3:1 ip-prefix 10.111.112.0/24	10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.254.4:1 ip-prefix 10.111.112.0/24	10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.254.4:1 ip-prefix 10.111.112.0/24	10.111.253.4	-	100	0	65100 65
* >	RD: 10.111.254.1:1 ip-prefix 10.111.134.0/24	-	-	-	0	i
* >Ec	RD: 10.111.254.2:1 ip-prefix 10.111.134.0/24	10.111.253.2	-	100	0	65100 65
* ec	RD: 10.111.254.2:1 ip-prefix 10.111.134.0/24	10.111.253.2	-	100	0	65100 65
* >Ec	RD: 10.111.254.3:1 ip-prefix 10.111.134.0/24	10.111.253.3	-	100	0	65100 65
* ec	RD: 10.111.254.3:1 ip-prefix 10.111.134.0/24	10.111.253.3	-	100	0	65100 65
* >Ec	RD: 10.111.254.4:1 ip-prefix 10.111.134.0/24	10.111.253.4	-	100	0	65100 65
* ec	RD: 10.111.254.4:1 ip-prefix 10.111.134.0/24	10.111.253.4	-	100	0	65100 65

```
s1-leaf1#show bgp evpn route-type ip-prefix ipv4 rd 10.111.254.4:1 detail
BGP routing table information for VRF default
Router identifier 10.111.254.1, local AS number 65101
BGP routing table entry for ip-prefix 10.111.112.0/24, Route Distinguisher: 10.111.254.
Paths: 2 available
65100 65104
  10.111.253.4 from 10.111.0.2 (10.111.0.2)
    Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
    Extended Community: Route-Target-AS:5001:5001 TunnelEncap:tunnelTypeVxlan EvpnRou
    VNI: 5001
65100 65104
  10.111.253.4 from 10.111.0.1 (10.111.0.1)
    Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
    Extended Community: Route-Target-AS:5001:5001 TunnelEncap:tunnelTypeVxlan EvpnRou
    VNI: 5001
```

```

BGP routing table entry for ip-prefix 10.111.134.0/24, Route Distinguisher: 10.111.254.
Paths: 2 available
 65100 65104
    10.111.253.4 from 10.111.0.2 (10.111.0.2)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP head, ECMP,
      Extended Community: Route-Target-AS:5001:5001 TunnelEncap:tunnelTypeVxlan EvpnRou
      VNI: 5001
 65100 65104
    10.111.253.4 from 10.111.0.1 (10.111.0.1)
      Origin IGP, metric -, localpref 100, weight 0, valid, external, ECMP, ECMP contri
      Extended Community: Route-Target-AS:5001:5001 TunnelEncap:tunnelTypeVxlan EvpnRou
      VNI: 5001

```

g. On **s1-leaf1**, check the local ARP and MAC address-table.

### Note

The MAC addresses in your lab may differ as they are randomly generated during the lab build. We see here that the ARP and MAC entry of **s1-host2** has been learned and imported via the Vxlan1 interface on **s1-leaf1** in both Host VLANs.

We also see the remote MAC of each VTEPs System ID including the highlighted one for **s1-leaf4** associated with VLAN 4093 and the Vxlan1 interface. This is how the local VTEP knows where to send routed (ie inter-subnet) traffic when destined to the remote MLAG pair. We can see this VLAN is dynamically created in the VLAN database and is mapped to our Layer 3 VNI (5001) in our VXLAN interface output. Be aware that since this VLAN is dynamic, the ID used in your lab may be different.

```
s1-leaf1#show ip arp vrf TENANT
```

Address	Age (sec)	Hardware Addr	Interface
10.111.112.201	0:05:14	001c.73c0.c616	Vlan112, Port-Channel5
10.111.112.202	-	001c.73c0.c617	Vlan112, Vxlan1
10.111.134.201	0:04:14	001c.73c0.c616	Vlan134, Port-Channel5
10.111.134.202	-	001c.73c0.c617	Vlan134, Vxlan1

```
s1-leaf1#show mac address-table dynamic
```

#### Mac Address Table

Vlan	Mac Address	Type	Ports	Moves	Last Move
112	001c.73c0.c616	DYNAMIC	Po5	1	0:05:27 ago
112	001c.73c0.c617	DYNAMIC	Vx1	1	0:04:15 ago
134	001c.73c0.c616	DYNAMIC	Po5	1	0:04:27 ago
134	001c.73c0.c617	DYNAMIC	Vx1	1	0:05:30 ago
4093	001c.73c0.c613	DYNAMIC	Vx1	1	1:00:13 ago
4093	001c.73c0.c614	DYNAMIC	Vx1	1	1:00:06 ago
4093	001c.73c0.c615	DYNAMIC	Vx1	1	0:52:35 ago

Total Mac Addresses for this criterion: 7

#### Multicast Mac Address Table

Vlan	Mac Address	Type	Ports
------	-------------	------	-------

Total Mac Addresses for this criterion: 0

```
s1-leaf1#show vlan 4093
```

VLAN	Name	Status	Ports
4093*	VLAN4093	active	Cpu, Vx1

\* indicates a Dynamic VLAN

```
s1-leaf1#show interfaces Vxlan1
```

Vxlan1 is up, line protocol is up (connected)

Hardware is Vxlan

Source interface is Loopback1 and is active with 10.111.253.1

Replication/Flood Mode is headend with Flood List Source: EVPN

Remote MAC learning via EVPN

VNI mapping to VLANs

Static VLAN to VNI mapping is

[112, 112] [134, 134]

Dynamic VLAN to VNI mapping for 'evpn' is

[4093, 5001]

Note: All Dynamic VLANs used by VCS are internal VLANs.

Use 'show vxlan vni' for details.

Static VRF to VNI mapping is

[TENANT, 5001]

Headend replication flood vtep list is:

112 10.111.253.3 10.111.253.4 10.111.253.2

```
134 10.111.253.3    10.111.253.4    10.111.253.2
Shared Router MAC is 0000.0000.0000
```

- h. On **s1-leaf1**, check the VXLAN data-plane for MAC address.

#### Note

Recall above that the Type 2 EVPN route for **s1-host2** was associated with an **ESI** and our Type 1 EVPN routes showed us that **s1-leaf3** and **s1-leaf4** are both members of that **ES**. Therefore we see two possible destination for this host MAC. The `show l2rib output mac <MAC of remote host>` command then allows us to see the VTEP info in the hardware showing us the load-balancing that will occur. Finally we can verify the ECMP path to the remote VTEP **s1-leaf4** via **s1-spine1** and **s1-spine2** with a simple `show ip route 10.111.253.4` command.



```
s1-leaf1#show vxlan address-table evpn
      Vxlan Mac Address Table
-----
```

VLAN	Mac Address	Type	Prt	VTEP	Moves	Last Move
112	001c.73c0.c617	EVPN	Vx1	10.111.253.3 10.111.253.4	1	0:07:51 ago
134	001c.73c0.c617	EVPN	Vx1	10.111.253.3 10.111.253.4	1	0:09:06 ago
4093	001c.73c0.c613	EVPN	Vx1	10.111.253.2	1	1:03:50 ago
4093	001c.73c0.c614	EVPN	Vx1	10.111.253.3	1	1:03:43 ago
4093	001c.73c0.c615	EVPN	Vx1	10.111.253.4	1	0:56:11 ago

```
Total Remote Mac Addresses for this criterion: 5
s1-leaf1#show l2rib output mac 001c.73c0.c617
001c.73c0.c617, VLAN 112, seq 1, pref 16, evpnDynamicRemoteMac, source: BGP
  Load Balance entry: 2-way
    VTEP 10.111.253.3
    VTEP 10.111.253.4
001c.73c0.c617, VLAN 134, seq 1, pref 16, evpnDynamicRemoteMac, source: BGP
  Load Balance entry: 2-way
    VTEP 10.111.253.3
    VTEP 10.111.253.4
s1-leaf1#show ip route 10.111.253.4

VRF: default
Codes: C - connected, S - static, K - kernel,
       O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
       E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
       N2 - OSPF NSSA external type2, B - Other BGP Routes,
       B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
       I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
       A O - OSPF Summary, NG - Nexthop Group Static Route,
       V - VXLAN Control Service, M - Martian,
       DH - DHCP client installed default route,
       DP - Dynamic Policy Route, L - VRF Leaked,
       G - gRIBI, RC - Route Cache Route

B E      10.111.253.4/32 [200/0] via 10.111.1.0, Ethernet2
                               via 10.111.2.0, Ethernet3
```

- i. On **s1-leaf1**, verify the Tenant Route table to ensure the Tenant networks on **s1-leaf4** has been installed in the overlay.

### Note

Note on the route table for the TENANT VRF, we see a single route entry for the tenant subnets since they are both locally attached.

Also note that the Type 2 MAC-IP Routes, which correspond to the ARP entry of **s1-host2** have also been installed as /32 host routes. This ensures that in a distributed VXLAN fabric, Layer 3 routed traffic is always directed to the VTEP where the host currently resides. This route is di-

rected to the shared MLAG VTEP IP and EVPN Router MAC. It will be ECMPed via the Spines providing a dual path for load-balancing and redundancy.

And again due to our Type 1 EVPN Routes, each /32 host is known to be attached to both **s1-leaf3** and **s1-leaf4** as they are members of the associated **ES**.

```
s1-leaf1#show ip route vrf TENANT
```

```
VRF: TENANT
```

```
Codes: C - connected, S - static, K - kernel,
        O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
        E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
        N2 - OSPF NSSA external type2, B - Other BGP Routes,
        B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
        I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
        A O - OSPF Summary, NG - Nexthop Group Static Route,
        V - VXLAN Control Service, M - Martian,
        DH - DHCP client installed default route,
        DP - Dynamic Policy Route, L - VRF Leaked,
        G - gRIBI, RC - Route Cache Route
```

```
Gateway of last resort is not set
```

```

B E      10.111.112.202/32 [200/0] via VTEP 10.111.253.3 VNI 5001 router-mac 00:1c:73:
          via VTEP 10.111.253.4 VNI 5001 router-mac 00:1c:73:
C        10.111.112.0/24 is directly connected, Vlan112
B E      10.111.134.202/32 [200/0] via VTEP 10.111.253.3 VNI 5001 router-mac 00:1c:73:
          via VTEP 10.111.253.4 VNI 5001 router-mac 00:1c:73:
C        10.111.134.0/24 is directly connected, Vlan134
```

## LAB COMPLETE!

© Copyright 2024, Arista Networks.

[Back to top](#)

Created using Sphinx (<http://sphinx-doc.org/>) 6.1.3.