


Layer 3 Leaf-Spine



(_images/nested_l3ls_topo_1.png)

Note

The manually-entered commands below that are part of this lab are equivalent to `L3LS_s1-leaf4_complete`.

1. Log into the **LabAccess** jumpserver:
 - a. Type `l3ls` at the prompt. The script will configure the datacenter with the exception of **s1-leaf4**.

Note

Did you know the "l3ls" script is composed of Python code that uses the CloudVision Portal REST API to automate the provisioning of CVP Configlets. The configlets that are configured via the REST API are `L3LS_s1-spine1`, `L3LS_s1-spine2`, `L3LS_s1-leaf1`, `L3LS_s1-leaf2`, `L3LS_s1-leaf3`, `L3LS_s1-leaf4`.

2. Configure SVI and VARP Virtual IP on the **s1-leaf4** switch using the following criteria
 - a. Create the vARP MAC Address in Global Configuration mode

Note

Arista EOS utilizes the Industry-Standard CLI. When entering configuration commands, be sure to first type `configure` to enter configuration mode.

```
ip virtual-router mac-address 00:1c:73:00:00:34
```

b. Create the VLAN, SVI and the Virtual Router Address

```
vlan 134
  name Host_Network_134
!
interface vlan 134
  ip address 10.111.134.3/24
  ip virtual-router address 10.111.134.1
```

c. Validate the configuration with the following:

```
s1-leaf4#show ip interface brief
```

| Interface | IP Address | Status | Protocol | MTU | Address Owner |
|-------------|-----------------|--------|----------|------|---------------|
| Management0 | 192.168.0.15/24 | up | up | 1500 | |
| Vlan134 | 10.111.134.3/24 | up | up | 1500 | |
| Vlan4094 | 10.255.255.2/30 | up | up | 1500 | |

```
s1-leaf4#show ip virtual-router
```

```
IP virtual router is configured with MAC address: 001c.7300.0034
```

```
IP virtual router address subnet routes not enabled
```

```
MAC address advertisement interval: 30 seconds
```

```
Protocol: U - Up, D - Down, T - Testing, UN - Unknown
```

```
NP - Not Present, LLD - Lower Layer Down
```

| Interface | Vrf | Virtual IP Address | Protocol | State |
|-----------|---------|--------------------|----------|--------|
| Vl134 | default | 10.111.134.1 | U | active |

3. Configure BGP on the **s1-leaf4** switch using the following criteria

a. Based on the diagram, configure L3 interfaces to **s1-spine1/s1-spine2** and interface Loopback0

```
interface Ethernet2
  description L3 Uplink - s1-spine1
  no switchport
  ip address 10.111.1.7/31
!
interface Ethernet3
  description L3 Uplink - s1-spine2
  no switchport
  ip address 10.111.2.7/31
!
interface Loopback0
  description Management and Router-id
  ip address 10.111.254.4/32
```

b. Validate the configuration with the following:

```
s1-leaf4#show ip interface brief
```

| Interface | IP Address | Status | Protocol | MTU | Address Owner |
|-------------|-----------------|--------|----------|-------|---------------|
| Ethernet2 | 10.111.1.7/31 | up | up | 1500 | |
| Ethernet3 | 10.111.2.7/31 | up | up | 1500 | |
| Loopback0 | 10.111.254.4/32 | up | up | 65535 | |
| Management0 | 192.168.0.15/24 | up | up | 1500 | |
| Vlan134 | 10.111.134.3/24 | up | up | 1500 | |
| Vlan4094 | 10.255.255.2/30 | up | up | 1500 | |

c. Based on the diagram, turn on BGP and configure the neighbor relationships on **s1-leaf4**. eBGP to **s1-spine1/s1-spine2** and iBGP to **s1-leaf3**.

Note

We are using a peer group to configure the neighbor attributes for the spines. This allows us to apply all bgp attributes within a group to each neighbor that is a member in a scalable method.

```
router bgp 65102
  router-id 10.111.254.4
  neighbor SPINE peer group
  neighbor SPINE remote-as 65100
  neighbor SPINE send-community standard extended
  neighbor 10.111.1.6 peer group SPINE
  neighbor 10.111.2.6 peer group SPINE
  neighbor 10.255.255.1 remote-as 65102
  neighbor 10.255.255.1 next-hop-self
```

Note

Since `neighbor 10.255.255.1 remote-as 65102` specifies an iBGP peering relationship (because the ASN is the same as this switch 65102), the receiving switch may not have a route to networks more than 1 hop away, hence the switches should each advertise that they are the next hop via the `neighbor 10.255.255.1 next-hop-self` statement. While this scenario is only 2 iBGP peers, in a network fabric with several iBGP peers, a switch inside an AS (and not on an edge) may not have a route to a switch in any external AS.

d. Validate the configuration and neighbor establishment

```
s1-leaf4(config-router-bgp)#show active
router bgp 65102
  router-id 10.111.254.4
  neighbor SPINE peer group
  neighbor SPINE remote-as 65100
  neighbor SPINE send-community standard extended
  neighbor 10.111.1.6 peer group SPINE
  neighbor 10.111.2.6 peer group SPINE
  neighbor 10.255.255.1 remote-as 65102
  neighbor 10.255.255.1 next-hop-self
s1-leaf4(config-router-bgp)#show ip bgp summary
BGP summary information for VRF default
Router identifier 10.111.254.4, local AS number 65102
Neighbor Status Codes: m - Under maintenance
```

| Neighbor | V AS | MsgRcvd | MsgSent | InQ | OutQ | Up/Down | State | PfxRcd | PfxAcc |
|--------------|---------|---------|---------|-----|------|----------|-------|--------|--------|
| 10.111.1.6 | 4 65100 | 10 | 8 | 0 | 0 | 00:01:02 | Estab | 5 | 5 |
| 10.111.2.6 | 4 65100 | 9 | 9 | 0 | 0 | 00:01:02 | Estab | 5 | 5 |
| 10.255.255.1 | 4 65102 | 9 | 8 | 0 | 0 | 00:01:00 | Estab | 9 | 9 |

4. Configure networks on **s1-leaf4** to advertise to **s1-spine1/s1-spine2**

- a. Add the following networks to BGP announcements on **s1-leaf4**:

```
router bgp 65102
  network 10.111.134.0/24
  network 10.111.254.4/32
```

- b. Verify that these networks are being advertised to the other Spines and Leafs

```
s1-leaf1#show ip route
```

```
VRF: default
```

```
Codes: C - connected, S - static, K - kernel,
```

```
0 - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
```

```
E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
```

```
N2 - OSPF NSSA external type2, B - Other BGP Routes,
```

```
B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
```

```
I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
```

```
A O - OSPF Summary, NG - Nexthop Group Static Route,
```

```
V - VXLAN Control Service, M - Martian,
```

```
DH - DHCP client installed default route,
```

```
DP - Dynamic Policy Route, L - VRF Leaked,
```

```
G - gRIBI, RC - Route Cache Route
```

```
Gateway of last resort is not set
```

```
B E    10.111.0.1/32 [200/0] via 10.111.1.0, Ethernet2
```

```
B E    10.111.0.2/32 [200/0] via 10.111.2.0, Ethernet3
```

```
C      10.111.1.0/31 is directly connected, Ethernet2
```

```
B E    10.111.1.0/24 [200/0] via 10.111.1.0, Ethernet2
```

```
C      10.111.2.0/31 is directly connected, Ethernet3
```

```
B E    10.111.2.0/24 [200/0] via 10.111.2.0, Ethernet3
```

```
C      10.111.112.0/24 is directly connected, Vlan112
```

```
B E    10.111.134.0/24 [200/0] via 10.111.1.0, Ethernet2
```

```
C      10.111.254.1/32 is directly connected, Loopback0
```

```
B I    10.111.254.2/32 [200/0] via 10.255.255.2, Vlan4094
```

```
B E    10.111.254.3/32 [200/0] via 10.111.1.0, Ethernet2
```

```
B E    10.111.254.4/32 [200/0] via 10.111.1.0, Ethernet2
```

```
C      10.255.255.0/30 is directly connected, Vlan4094
```

```
C      192.168.0.0/24 is directly connected, Management0
```

```
s1-leaf1#show ip bgp
```

```
BGP routing table information for VRF default
```

```
Router identifier 10.111.254.1, local AS number 65101
```

```
Route status codes: s - suppressed, * - valid, > - active, E - ECMP head, e - ECMP
```

```
S - Stale, c - Contributing to ECMP, b - backup, L - labeled-unicast
```

```
% - Pending BGP convergence
```

```
Origin codes: i - IGP, e - EGP, ? - incomplete
```

```
RPKI Origin Validation codes: V - valid, I - invalid, U - unknown
```

```
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link Local
```

| | Network | Next Hop | Metric | AIGP | LocPref | Wei |
|-----|---------------|--------------|--------|------|---------|-----|
| * > | 10.111.0.1/32 | 10.111.1.0 | 0 | - | 100 | 0 |
| * | 10.111.0.1/32 | 10.255.255.2 | 0 | - | 100 | 0 |
| * > | 10.111.0.2/32 | 10.111.2.0 | 0 | - | 100 | 0 |
| * | 10.111.0.2/32 | 10.255.255.2 | 0 | - | 100 | 0 |
| * > | 10.111.1.0/24 | 10.111.1.0 | 0 | - | 100 | 0 |
| * | 10.111.1.0/24 | 10.255.255.2 | 0 | - | 100 | 0 |
| * > | 10.111.2.0/24 | 10.111.2.0 | 0 | - | 100 | 0 |
| * | 10.111.2.0/24 | 10.255.255.2 | 0 | - | 100 | 0 |

```

* > 10.111.112.0/24 - - - - 0
* 10.111.112.0/24 10.255.255.2 0 - 100 0
* > 10.111.134.0/24 10.111.1.0 0 - 100 0
* 10.111.134.0/24 10.111.2.0 0 - 100 0
* 10.111.134.0/24 10.255.255.2 0 - 100 0
* > 10.111.254.1/32 - - - - 0
* > 10.111.254.2/32 10.255.255.2 0 - 100 0
* > 10.111.254.3/32 10.111.1.0 0 - 100 0
* 10.111.254.3/32 10.111.2.0 0 - 100 0
* 10.111.254.3/32 10.255.255.2 0 - 100 0
* > 10.111.254.4/32 10.111.1.0 0 - 100 0
* 10.111.254.4/32 10.111.2.0 0 - 100 0
* 10.111.254.4/32 10.255.255.2 0 - 100 0

```

```
s1-leaf1#show ip route bgp
```

```
VRF: default
```

```

Codes: C - connected, S - static, K - kernel,
       O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
       E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
       N2 - OSPF NSSA external type2, B - Other BGP Routes,
       B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
       I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
       A O - OSPF Summary, NG - Nexthop Group Static Route,
       V - VXLAN Control Service, M - Martian,
       DH - DHCP client installed default route,
       DP - Dynamic Policy Route, L - VRF Leaked,
       G - gRIBI, RC - Route Cache Route

```

```

B E 10.111.0.1/32 [200/0] via 10.111.1.0, Ethernet2
B E 10.111.0.2/32 [200/0] via 10.111.2.0, Ethernet3
B E 10.111.1.0/24 [200/0] via 10.111.1.0, Ethernet2
B E 10.111.2.0/24 [200/0] via 10.111.2.0, Ethernet3
B E 10.111.134.0/24 [200/0] via 10.111.1.0, Ethernet2
B I 10.111.254.2/32 [200/0] via 10.255.255.2, Vlan4094
B E 10.111.254.3/32 [200/0] via 10.111.1.0, Ethernet2
B E 10.111.254.4/32 [200/0] via 10.111.1.0, Ethernet2

```

- c. Add in multiple paths by enabling ECMP, on **s1-leaf4**, jump into BGP configuration mode and add:

```

router bgp 65102
  maximum-paths 2

```

- d. Check the BGP and IP route tables on **s1-leaf4** as well as each of the **Spines** and **Leafs**

```
s1-spine1#show ip route
```

```
VRF: default
```

```
Codes: C - connected, S - static, K - kernel,
```

```
0 - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
```

```
E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
```

```
N2 - OSPF NSSA external type2, B - Other BGP Routes,
```

```
B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
```

```
I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
```

```
A O - OSPF Summary, NG - Nexthop Group Static Route,
```

```
V - VXLAN Control Service, M - Martian,
```

```
DH - DHCP client installed default route,
```

```
DP - Dynamic Policy Route, L - VRF Leaked,
```

```
G - gRIBI, RC - Route Cache Route
```

```
Gateway of last resort is not set
```

```
C      10.111.0.1/32 is directly connected, Loopback0
C      10.111.1.0/31 is directly connected, Ethernet2
C      10.111.1.2/31 is directly connected, Ethernet3
C      10.111.1.4/31 is directly connected, Ethernet4
C      10.111.1.6/31 is directly connected, Ethernet5
S      10.111.1.0/24 is directly connected, Null0
B E    10.111.112.0/24 [200/0] via 10.111.1.1, Ethernet2
                        via 10.111.1.3, Ethernet3
B E    10.111.134.0/24 [200/0] via 10.111.1.5, Ethernet4
                        via 10.111.1.7, Ethernet5
B E    10.111.254.1/32 [200/0] via 10.111.1.1, Ethernet2
                        via 10.111.1.3, Ethernet3
B E    10.111.254.2/32 [200/0] via 10.111.1.1, Ethernet2
                        via 10.111.1.3, Ethernet3
B E    10.111.254.3/32 [200/0] via 10.111.1.5, Ethernet4
                        via 10.111.1.7, Ethernet5
B E    10.111.254.4/32 [200/0] via 10.111.1.5, Ethernet4
                        via 10.111.1.7, Ethernet5
C      192.168.0.0/24 is directly connected, Management0
```

```
s1-spine1#show ip bgp
```

```
BGP routing table information for VRF default
```

```
Router identifier 10.111.0.1, local AS number 65100
```

```
Route status codes: s - suppressed, * - valid, > - active, E - ECMP head, e - ECMP
                    S - Stale, c - Contributing to ECMP, b - backup, L - labeled-unicast
                    % - Pending BGP convergence
```

```
Origin codes: i - IGP, e - EGP, ? - incomplete
```

```
RPKI Origin Validation codes: V - valid, I - invalid, U - unknown
```

```
AS Path Attributes: Or-ID - Originator ID, C-LST - Cluster List, LL Nexthop - Link List
```

| | Network | Next Hop | Metric | AIGP | LocPref | Weight |
|-------|-----------------|------------|--------|------|---------|--------|
| * > | 10.111.0.1/32 | - | - | - | - | 0 |
| * > | 10.111.1.0/24 | - | - | - | - | 0 |
| * >Ec | 10.111.112.0/24 | 10.111.1.1 | 0 | - | 100 | 0 |
| * ec | 10.111.112.0/24 | 10.111.1.3 | 0 | - | 100 | 0 |

```
* >Ec 10.111.134.0/24 10.111.1.5 0 - 100 0
* ec 10.111.134.0/24 10.111.1.7 0 - 100 0
* >Ec 10.111.254.1/32 10.111.1.1 0 - 100 0
* ec 10.111.254.1/32 10.111.1.3 0 - 100 0
* >Ec 10.111.254.2/32 10.111.1.3 0 - 100 0
* ec 10.111.254.2/32 10.111.1.1 0 - 100 0
* >Ec 10.111.254.3/32 10.111.1.5 0 - 100 0
* ec 10.111.254.3/32 10.111.1.7 0 - 100 0
* >Ec 10.111.254.4/32 10.111.1.7 0 - 100 0
* ec 10.111.254.4/32 10.111.1.5 0 - 100 0
```

```
s1-spine1#sh ip route bgp
```

```
VRF: default
```

```
Codes: C - connected, S - static, K - kernel,
        O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
        E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
        N2 - OSPF NSSA external type2, B - Other BGP Routes,
        B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
        I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
        A O - OSPF Summary, NG - Nexthop Group Static Route,
        V - VXLAN Control Service, M - Martian,
        DH - DHCP client installed default route,
        DP - Dynamic Policy Route, L - VRF Leaked,
        G - gRIBI, RC - Route Cache Route
```

```
B E 10.111.112.0/24 [200/0] via 10.111.1.1, Ethernet2
                        via 10.111.1.3, Ethernet3
B E 10.111.134.0/24 [200/0] via 10.111.1.5, Ethernet4
                        via 10.111.1.7, Ethernet5
B E 10.111.254.1/32 [200/0] via 10.111.1.1, Ethernet2
                        via 10.111.1.3, Ethernet3
B E 10.111.254.2/32 [200/0] via 10.111.1.1, Ethernet2
                        via 10.111.1.3, Ethernet3
B E 10.111.254.3/32 [200/0] via 10.111.1.5, Ethernet4
                        via 10.111.1.7, Ethernet5
B E 10.111.254.4/32 [200/0] via 10.111.1.5, Ethernet4
                        via 10.111.1.7, Ethernet5
```

Note

ECMP is now working - notice the new status code in the *show ip bgp* output on s1-leaf4

5. Validate connectivity from **s1-host1** to **s1-host2**. From **s1-host1** execute:

```
ping 10.111.134.202
traceroute 10.111.134.202
```

- Verify **s1-leaf4**'s IP address is in the traceroute path, either interface 10.111.1.7 via **s1-spine1** or interface 10.111.2.7 via **s1-spine2**. If traffic is hashing via **s1-leaf3**'s 10.111.1.5 or 10.111.2.5 interfaces perform the optional `shutdown` steps below on **s1-leaf3**


```
router bgp 65102
  neighbor 10.111.1.4 shutdown
  neighbor 10.111.2.4 shutdown
```

- b. Rerun traceroute/verification from **s1-host1** to **s1-host2** then revert the shutdown changes on **s1-leaf3**

```
router bgp 65102
  no neighbor 10.111.1.4 shutdown
  no neighbor 10.111.2.4 shutdown
```

6. Other BGP features to play with if you have time:

- a. Route Redistribution: For fun, do a `watch 1 diff show ip route | begin Gateway` on **s1-leaf1** & **s1-leaf2** and let those run while you execute the command `redistribute connected` below on **s1-leaf3**. You will see new routes being injected into the route tables of **s1-leaf1** & **s1-leaf2**.

```
router bgp 65102
  redistribute connected
```

- b. Route Maps and Prefix-Lists: Below is an example of some basic Prefix-Lists and Route-Maps that can be used for BGP filtering. Note that this is just an example and will not impact route advertisement in the lab.

```
<Example>

ip prefix-list BOGON-Prefixes seq 10 permit 10.0.0.0/8
ip prefix-list BOGON-Prefixes seq 20 permit 172.16.0.0/12
ip prefix-list BOGON-Prefixes seq 30 permit 192.168.0.0/16
!
route-map BOGONS permit 10
  match ip address prefix-list BOGON-Prefixes
!
route-map BOGONS deny 20
!
route-map InboundSP1 deny 10
  sub-route-map BOGONS
!
route-map InboundSP1 permit 20
  set local-preference 200
!
router bgp 65102
  neighbor UpstreamSP1 route-map InboundSP1 in
```

- c. BFD: BFD is a low-overhead, protocol-independent mechanism which adjacent systems can use instead for faster detection of faults in the path between them. BFD is a simple mechanism which detects the liveness of a connection between adjacent systems, allowing it to quickly detect failure of any element in the connection. Note that BFD is not running on the other devices so the BFD neighbor will not come up.

```
router bgp 65102
  neighbor <neighbor_ip> bfd
```

7. Troubleshooting BGP:

```
show ip bgp summary
show ip bgp
show ip bgp neighbor <neighbor_ip>
show run section bgp
show log
```

LAB COMPLETE!

© Copyright 2024, Arista Networks.
Created using Sphinx (<http://sphinx-doc.org/>) 6.1.3.

[Back to top](#)