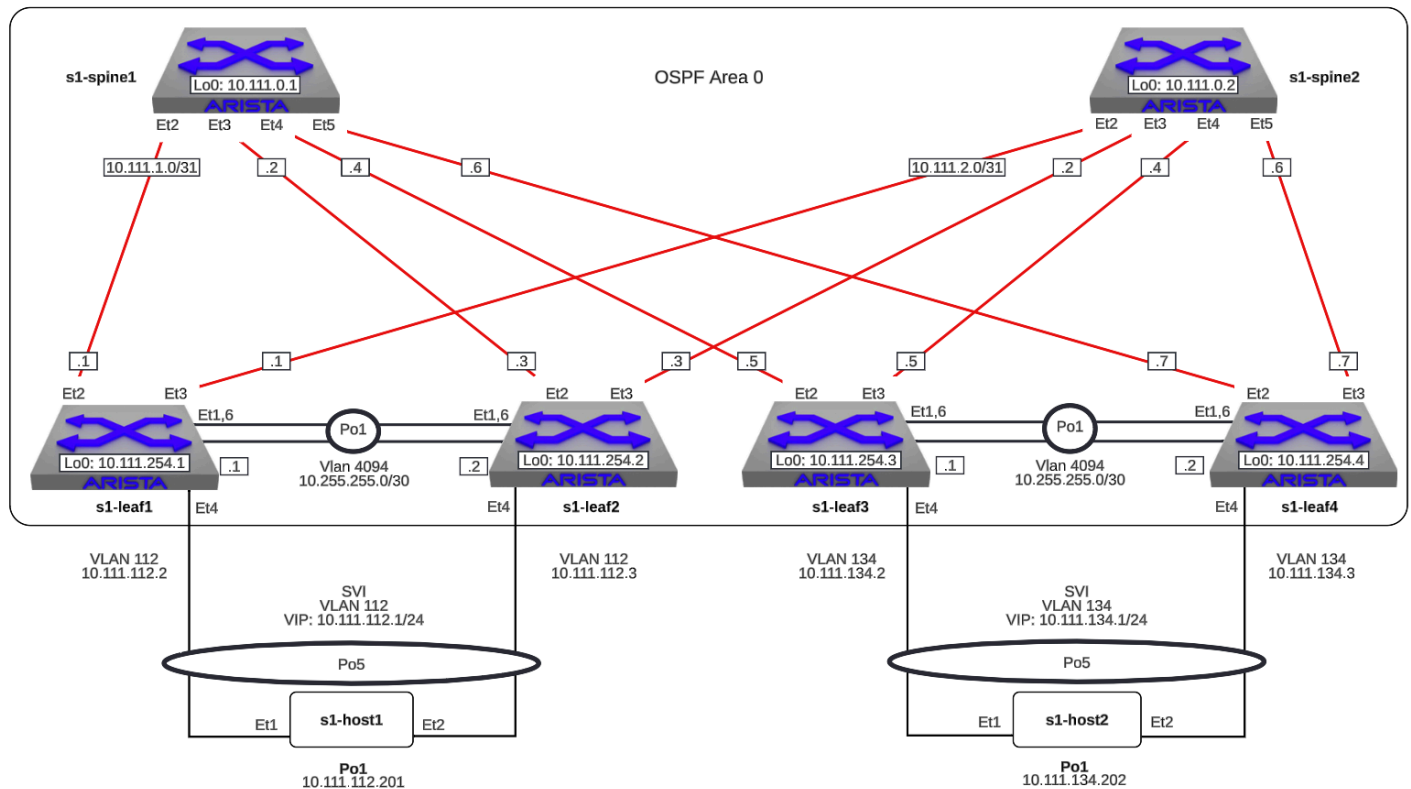


Layer 3 Leaf-Spine with OSPF



(_images/nested_l3ls-ospf_topo_1.png)

Note

The manually-entered commands below that are part of this lab are equivalent to `L3LS-OSPF_s1-leaf4_complete`.

1. Log into the **LabAccess** jumpserver:
 - a. Type `l3ls-ospf` at the prompt. The script will configure the datacenter with the exception of **s1-leaf4**.

Note

Did you know the "l3ls-ospf" script is composed of Python code that uses the CloudVision Portal REST API to automate the provisioning of CVP Configlets. The configlets that are configured via the REST API are `L3LS-OSPF_s1-spine1`, `L3LS-OSPF_s1-spine2`, `L3LS-OSPF_s1-leaf1`, `L3LS-OSPF_s1-leaf2`, `L3LS-OSPF_s1-leaf3`, `L3LS-OSPF_s1-leaf4`.

2. Configure SVI and VARP Virtual IP on the **s1-leaf4** switch using the following criteria
 - a. Create the vARP MAC Address in Global Configuration mode

Note

Arista EOS utilizes the Industry-Standard CLI. When entering configuration commands, be sure to first type `configure` to enter configuration mode.

```
ip virtual-router mac-address 00:1c:73:00:00:34
```

b. Create the VLAN, SVI and the Virtual Router Address

```
vlan 134
  name Host_Network_134
!
interface vlan 134
  ip address 10.111.134.3/24
  ip virtual-router address 10.111.134.1
```

c. Validate the configuration with the following:

```
s1-leaf4#show ip interface brief
```

Interface	IP Address	Status	Protocol	MTU	Address Owner
Management0	192.168.0.15/24	up	up	1500	
Vlan134	10.111.134.3/24	up	up	1500	
Vlan4094	10.255.255.2/30	up	up	1500	

```
s1-leaf4#show ip virtual-router
```

```
IP virtual router is configured with MAC address: 001c.7300.0034
IP virtual router address subnet routes not enabled
MAC address advertisement interval: 30 seconds
```

```
Protocol: U - Up, D - Down, T - Testing, UN - Unknown
          NP - Not Present, LLD - Lower Layer Down
```

Interface	Vrf	Virtual IP Address	Protocol	State
Vl134	default	10.111.134.1	U	active

3. Configure OSPF on the **s1-leaf4** switch using the following criteria

a. Based on the diagram, configure L3 interfaces to **s1-spine1/s1-spine2** and interface Loopback0

```

interface Ethernet2
  description L3 Uplink - s1-spine1
  no switchport
  ip address 10.111.1.7/31
!
interface Ethernet3
  description L3 Uplink - s1-spine2
  no switchport
  ip address 10.111.2.7/31
!
interface Loopback0
  description Management and Router-id
  ip address 10.111.254.4/32

```

b. Validate the configuration with the following:

```
s1-leaf4#show ip interface brief
```

Interface	IP Address	Status	Protocol	MTU	Address
					Owner
Ethernet2	10.111.1.7/31	up	up	1500	
Ethernet3	10.111.2.7/31	up	up	1500	
Loopback0	10.111.254.4/32	up	up	65535	
Management0	192.168.0.15/24	up	up	1500	
Vlan134	10.111.134.3/24	up	up	1500	
Vlan4094	10.255.255.2/30	up	up	1500	

c. Based on the diagram, enable OSPF and configure the interfaces on **s1-leaf4**. Connections to **s1-spine1/s1-spine2** and **s1-leaf3** will be part of Area 0.

Note

In EOS, process-level configuration happens with the OSPF context and peer-specific configuration such as Area and Authentication happen under the interface.

```

interface Ethernet2
  ip ospf area 0.0.0.0
  ip ospf network point-to-point
!
interface Ethernet3
  ip ospf area 0.0.0.0
  ip ospf network point-to-point
!
interface Vlan4094
  ip ospf area 0.0.0.0
  ip ospf network point-to-point
!
router ospf 100
  router-id 10.111.254.4

```

Note

We are leveraging OSPF Point-to-Point networks to eliminate the need for DR elections on non-broadcast networks.

- d. Validate the configuration and process status.

```
s1-leaf4(config-router-ospf)#show active
router ospf 100
  router-id 10.111.254.4
  max-lsa 12000
s1-leaf4(config-router-ospf)#show run interfaces Ethernet 2-3
interface Ethernet2
  description L3 Uplink - s1-spine1
  no switchport
  ip address 10.111.1.7/31
  ip ospf area 0.0.0.0
  ip ospf network point-to-point
interface Ethernet3
  description L3 Uplink - s1-spine2
  no switchport
  ip address 10.111.2.7/31
  ip ospf area 0.0.0.0
  ip ospf network point-to-point
s1-leaf4(config-router-ospf)#show ip ospf summary
OSPF instance 100 with ID 10.111.254.4, VRF default
Time since last SPF: 122 s
Max LSAs: 12000, Total LSAs: 1
Type-5 Ext LSAs: 0
```

ID	Type	Intf	Nbrs (full)	RTR LSA	NW LSA	SUM LSA	ASBR LSA	TYPE-7 LSA
0.0.0.0	normal	3	0 (0)	1	0	0	0	0

4. Enable OSPF Authentication on **s1-leaf4** to peer to **s1-spine1/s1-spine2** and **s1-leaf3**

- a. Add the following Authentication commands to OSPF Interfaces on **s1-leaf4**:

```
interface Ethernet2
  ip ospf authentication message-digest
  ip ospf message-digest-key 1 sha512 Arista123!
!
interface Ethernet3
  ip ospf authentication message-digest
  ip ospf message-digest-key 1 sha512 Arista123!
!
interface Vlan4094
  ip ospf authentication message-digest
  ip ospf message-digest-key 1 sha512 Arista123!
```

- b. Verify that peering is established to directly connected neighbors.

```
s1-leaf4(config-if-Vl4094)#show ip ospf neighbor
```

Neighbor ID	Instance	VRF	Pri	State	Dead Time	Address
10.111.0.1	100	default	0	FULL	00:00:38	10.111.1.6
10.111.0.2	100	default	0	FULL	00:00:37	10.111.2.6
10.111.254.3	100	default	1	FULL	00:00:31	10.255.255.1

- c. Enable OSPF Advertisement of local networks on **s1-leaf4**.

```
interface Loopback0
  ip ospf area 0.0.0.0
!
interface Vlan134
  ip ospf area 0.0.0.0
```

- d. However, we do not want to form OSPF adjacencies on these interfaces, so enable passive-interface functionality on **s1-leaf4**.

```
router ospf 100
  passive-interface default
  no passive-interface Ethernet2
  no passive-interface Ethernet3
  no passive-interface Vlan4094
```

- e. Check the OSPF Database and IP route tables on **s1-leaf4** as well as each of the **Spines** and **Leafs**

```
s1-leaf4#show ip ospf database
```

```
OSPF Router with ID(10.111.254.4) (Instance ID 100) (VRF default)
```

```
Router Link States (Area 0.0.0.0)
```

Link ID	ADV Router	Age	Seq#	Checksum	Link count
10.111.0.2	10.111.0.2	356	0x80000011	0x4670	6
10.111.0.1	10.111.0.1	355	0x80000011	0x3be	6
10.111.254.3	10.111.254.3	358	0x80000013	0x691	6
10.111.254.2	10.111.254.2	1198	0x8000000f	0x58f3	5
10.111.254.4	10.111.254.4	354	0x80000018	0x3b28	8
10.111.254.1	10.111.254.1	1198	0x8000000f	0x5ff8	5

```
Network Link States (Area 0.0.0.0)
```

Link ID	ADV Router	Age	Seq#	Checksum
10.255.255.2	10.111.254.2	3600	0x80000116	0x6448
10.111.2.3	10.111.254.2	1198	0x80000002	0x5ded
10.111.2.1	10.111.254.1	1198	0x80000002	0x6de1
10.111.2.5	10.111.254.3	1198	0x80000002	0x4df9
10.111.1.1	10.111.254.1	1198	0x80000002	0x6ae6
10.111.1.5	10.111.254.3	1198	0x80000002	0x4afe
10.111.1.3	10.111.254.2	1198	0x80000002	0x5af2

```
s1-leaf4#show ip ospf database detail 10.111.254.1
```

```
OSPF Router with ID(10.111.254.4) (Instance ID 100) (VRF default)
```

```
LS Age: 1294
```

```
Options: (E DC)
```

```
LS Type: Router Links
```

```
Link State ID: 10.111.254.1
```

```
Advertising Router: 10.111.254.1
```

```
LS Seq Number: 0x8000000f
```

```
Checksum: 0x5ff8
```

```
Length: 84
```

```
Number of Links: 5
```

```
Link connected to: a Transit Network
```

```
(Link ID) Designated Router address: 10.111.1.1
```

```
(Link Data) Router Interface address: 10.111.1.1
```

```
Number of TOS metrics: 0
```

```
TOS 0 Metrics: 10
```

```
Link connected to: a Stub Network
```

```
(Link ID) Network/subnet number: 10.111.254.1
```

```
(Link Data) Network Mask: 255.255.255.255
```

```
Number of TOS metrics: 0
```

```
TOS 0 Metrics: 10
```

```

Link connected to: a Transit Network
(Link ID) Designated Router address: 10.111.2.1
(Link Data) Router Interface address: 10.111.2.1
Number of TOS metrics: 0
TOS 0 Metrics: 10

```

```

Link connected to: a Transit Network
(Link ID) Designated Router address: 10.255.255.2
(Link Data) Router Interface address: 10.255.255.1
Number of TOS metrics: 0
TOS 0 Metrics: 10

```

```

Link connected to: a Stub Network
(Link ID) Network/subnet number: 10.111.112.0
(Link Data) Network Mask: 255.255.255.0
Number of TOS metrics: 0
TOS 0 Metrics: 10

```

```
s1-leaf4#show ip route 10.111.112.0/24
```

```
VRF: default
```

```
Source Codes:
```

```

C - connected, S - static, K - kernel,
O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,
E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,
N2 - OSPF NSSA external type2, B - Other BGP Routes,
B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,
I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,
A O - OSPF Summary, NG - Nexthop Group Static Route,
V - VXLAN Control Service, M - Martian,
DH - DHCP client installed default route,
DP - Dynamic Policy Route, L - VRF Leaked,
G - gRIBI, RC - Route Cache Route,
CL - CBF Leaked Route

```

```

O          10.111.112.0/24 [110/30]
            via 10.111.1.6, Ethernet2
            via 10.111.2.6, Ethernet3

```

Note

ECMP is automatically enabled in OSPF as it is an IGP.

5. Validate connectivity from **s1-host1** to **s1-host2**. From **s1-host1** execute:

```

ping 10.111.134.202
traceroute 10.111.134.202

```

- a. Verify **s1-leaf4**'s IP address is in the traceroute path, either interface 10.111.1.7 via **s1-spine1** or interface 10.111.2.7 via **s1-spine2**. If traffic is hashing via **s1-leaf3**'s 10.111.1.5 or 10.111.2.5 interfaces perform the optional `shutdown` steps below on **s1-leaf3**

```
interface Ethernet2-3
shutdown
```

- b. Rerun traceroute/verification from **s1-host1** to **s1-host2** then revert the `shutdown` changes on **s1-leaf3**

```
interface Ethernet2-3
no shutdown
```

6. Other OSPF features to play with if you have time:

- a. Route Redistribution: For fun, do a `watch 1 diff show ip route | begin Gateway` on **s1-leaf1** and let that run while you execute the following commands on **s1-leaf4**. You will see the new route being injected into the route table of **s1-leaf1**.

```
ip route 0.0.0.0/0 Null0
!
router ospf 100
redistribute static
```

- b. BFD: BFD is a low-overhead, protocol-independent mechanism which adjacent systems can use instead for faster detection of faults in the path between them. BFD is a simple mechanism which detects the liveness of a connection between adjacent systems, allowing it to quickly detect failure of any element in the connection. Note that BFD is not running on the other devices so the BFD neighbor will not come up until you configure it on multiple devices.

```
router ospf 100
bfd default
```

7. Troubleshooting BGP:

```
show ip ospf summary
show ip ospf
show ip ospf neighbor <neighbor_ip>
show run section ospf
show log
```

LAB COMPLETE!