


# VxLAN



## Note

Did you know the `vxlan` script is composed of Python code that uses the CloudVision Portal REST API to automate the provisioning of CVP configlets. The configlets that are configured via the REST API are `VXLAN_s1-spine1`, `VXLAN_s1-spine2`, `VXLAN_s1-leaf1`, `VXLAN_s1-leaf2`, `VXLAN_s1-leaf3`, `VXLAN_s1-leaf4`. In addition each leaf also gets the `VLANs` configlet.

## Note

The manually-entered commands below that are part of this lab are equivalent to `VXLAN_s1-leaf4_complete`.

1. Log into the LabAccess jumpserver:
  1. Type `vxlan` at the prompt. The script will configure the datacenter with the exception of **s1-leaf4**.
2. On **s1-leaf4**, configure Port-channels connecting to **s1-host2**

```
configure
interface port-channel 5
    description MLAG - HOST2
    switchport access vlan 112
    mlag 5

interface Ethernet4
    description HOST2
    channel-group 5 mode active
    lacp timer fast
```

### 3. Verify MLAG on **s1-leaf4**

```
s1-leaf4(config)#show mlag
MLAG Configuration:
domain-id                :                MLAG
local-interface          :                Vlan4094
peer-address             :                10.255.255.1
peer-link                :                Port-Channel1
peer-config              :                inconsistent

MLAG Status:
state                    :                Active
negotiation status      :                Connected
peer-link status        :                Up
local-int status        :                Up
system-id               :                02:1c:73:c0:c6:14
dual-primary detection  :                Disabled
dual-primary interface errdisabled :      False

MLAG Ports:
Disabled                 :                0
Configured               :                0
Inactive                 :                0
Active-partial           :                0
Active-full              :                1

s1-leaf4(config)#show mlag interfaces

      mlag      desc      state      local      remote      local/remote
-----
      5      MLAG - HOST2      active-full      Po5      Po5      up/up

s1-leaf4(config)#show port-channel dense

      Flags
-----
a - LACP Active      p - LACP Passive      * - static fallback
F - Fallback enabled  f - Fallback configured  ^ - individual fallback
U - In Use           D - Down
+ - In-Sync          - - Out-of-Sync      i - incompatible with agg
P - bundled in Po    s - suspended          G - Aggregable
I - Individual       S - ShortTimeout       w - wait for agg
E - Inactive. The number of configured port channels exceeds the config limit
M - Exceeds maximum weight

Number of channels in use: 2
Number of aggregators: 2

      Port-Channel      Protocol      Ports
-----
      Po1(U)            LACP(a)      Et1(PG+) Et6(PG+)
```

Po5(U)

LACP(a)

Et4(PSG+) PEt4(P)

#### 4. Validate BGP operation **s1-leaf4**

```
s1-leaf4(config)#sh run sec bgp
router bgp 65102
  router-id 10.111.254.4
  maximum-paths 2
  neighbor SPINE peer group
  neighbor SPINE remote-as 65100
  neighbor SPINE send-community standard extended
  neighbor 10.111.1.6 peer group SPINE
  neighbor 10.111.2.6 peer group SPINE
  neighbor 10.255.255.1 remote-as 65102
  neighbor 10.255.255.1 next-hop-self
  network 10.111.112.0/24
  network 10.111.134.0/24
  network 10.111.254.4/32
```

```
s1-leaf4(config)#show ip route bgp
```

VRF: default

Codes: C - connected, S - static, K - kernel,  
 O - OSPF, IA - OSPF inter area, E1 - OSPF external type 1,  
 E2 - OSPF external type 2, N1 - OSPF NSSA external type 1,  
 N2 - OSPF NSSA external type2, B - Other BGP Routes,  
 B I - iBGP, B E - eBGP, R - RIP, I L1 - IS-IS level 1,  
 I L2 - IS-IS level 2, O3 - OSPFv3, A B - BGP Aggregate,  
 A O - OSPF Summary, NG - Nexthop Group Static Route,  
 V - VXLAN Control Service, M - Martian,  
 DH - DHCP client installed default route,  
 DP - Dynamic Policy Route, L - VRF Leaked,  
 G - gRIBI, RC - Route Cache Route

```
B E    10.111.0.1/32 [200/0] via 10.111.1.6, Ethernet2
B E    10.111.0.2/32 [200/0] via 10.111.2.6, Ethernet3
B E    10.111.1.0/24 [200/0] via 10.111.1.6, Ethernet2
B E    10.111.2.0/24 [200/0] via 10.111.2.6, Ethernet3
B E    10.111.253.1/32 [200/0] via 10.111.1.6, Ethernet2
      via 10.111.2.6, Ethernet3
B I    10.111.253.3/32 [200/0] via 10.255.255.1, Vlan4094
B E    10.111.254.1/32 [200/0] via 10.111.1.6, Ethernet2
      via 10.111.2.6, Ethernet3
B E    10.111.254.2/32 [200/0] via 10.111.1.6, Ethernet2
      via 10.111.2.6, Ethernet3
B I    10.111.254.3/32 [200/0] via 10.255.255.1, Vlan4094
```

```
s1-leaf4(config)#show ip interface brief
```

Interface	IP Address	Status	Protocol	MTU	Address Owner
Ethernet2	10.111.1.7/31	up	up	1500	
Ethernet3	10.111.2.7/31	up	up	1500	
Loopback0	10.111.254.4/32	up	up	65535	

Management0	192.168.0.15/24	up	up	1500
Vlan112	10.111.112.1/24	up	up	1500
Vlan134	10.111.134.1/24	up	up	1500
Vlan4094	10.255.255.2/30	up	up	1500

```
s1-leaf4(config)#show ip bgp summary
```

```
BGP summary information for VRF default
```

```
Router identifier 10.111.254.4, local AS number 65102
```

```
Neighbor Status Codes: m - Under maintenance
```

Neighbor	V AS	MsgRcvd	MsgSent	InQ	OutQ	Up/Down	State	PfxRcd	PfxA
10.111.1.6	4 65100	333	335	0	0	04:34:48	Estab	5	5
10.111.2.6	4 65100	329	332	0	0	04:34:58	Estab	6	6
10.255.255.1	4 65100	335	333	0	0	04:34:40	Estab	11	11

## Note

`show ip bgp summary` will show that the BGP neighbors have moved to `Estab` state. Note the iBGP peering between Leaf3 & Leaf4. Also note the route to the shared loopback1 of Leaf1 & Leaf2. This is the remote VTEP on the other side of the leaf-spine network.

## 5. Create Loopback 1 and the VXLAN VTEP (VTI) interfaces on **s1-leaf4**

### 1. Configuration

```
configure
interface Loopback1
  ip address 10.111.253.3/32

interface vxlan 1
  vxlan source-interface loopback 1
  vxlan vlan 112 vni 112
  vxlan flood vtep 10.111.253.1
```

## Note

`vxlan flood vtep 10.111.253.1` adds the shared loopback1 IP address on Leaf1 & Leaf2 to the HER list. Note that for autodiscovery of VTEPs, one must use BGP eVPN (see eVPN labs) or CVX (see CVX lab).

### 2. Verification

```
s1-leaf4(config)#show run int vxlan1
interface Vxlan1
  vxlan source-interface Loopback1
  vxlan udp-port 4789
  vxlan vlan 112 vni 112
  vxlan flood vtep 10.111.253.1

s1-leaf4(config-if-Vx1)#sh int vxlan 1
Vxlan1 is up, line protocol is up (connected)
Hardware is Vxlan
Source interface is Loopback1 and is active with 10.111.253.3
Replication/Flood Mode is headend with Flood List Source: CLI
Remote MAC learning via Datapath
VNI mapping to VLANs
Static VLAN to VNI mapping is
  [12, 112]
Note: All Dynamic VLANs used by VCS are internal VLANs.
      Use 'show vxlan vni' for details.
Static VRF to VNI mapping is not configured
Headend replication flood vtep list is:
  12 10.111.253.1
MLAG Shared Router MAC is 0000.0000.0000
```

6. Log into **s1-host1** and **s1-host2**, ping the vARP VIP and the other host

1. Host 1 ping tests. From **s1-host1**:

```
ping 10.111.112.1
ping 10.111.112.202
```

#### Note

The TTL in the ping outputs above. Even though .202 is many switches away, it appears locally connected and has the same TTL as the ping to .1. It's also interesting to realize that due to MLAG hashing of both the ARP requests and ping packet flows that pings to the SVI addresses of .2 & .3 may or may not work. Do you know why?

2. Host 1 MAC/ARP information

```
s1-host1#show interfaces po1 | grep -i Hardware
Hardware is Port-Channel, address is 001c.73c0.c616

s1-host1#show arp
Address          Age (sec)  Hardware Addr  Interface
192.168.0.1      0:00:00   124e.b1e1.7180  Management0
192.168.0.5      0:00:05   001c.73a0.c601  Management0
10.111.112.1     0:38:05   001c.7300.0001  Port-Channel1
10.111.112.202   0:14:05   001c.73c0.c617  Port-Channel1
```

**Note**

Note the MAC addresses returned by the commands above.

3. Host 2 ping tests. From **s1-host2**:

```
ping 10.111.112.1
ping 10.111.112.201
```

**Note**

Note the TTL in the ping outputs above. Even though .201 is many switches away, it appears locally connected and has the same TTL as the ping to .1. Also note that the vARP VIP (10.111.112.1) address & and vARP MAC address (00:1c:73:00:00:ff) are the **same** for both leaf pairs - this IP address is known as an AnyCast IP address. If a VM was motioning from **s1-host1** to **s1-host2** for maintenance, the default GW address nor the ARP cache on that VM need to change.

## 4. Host 2 MAC/ARP information

```
show interface po1 | grep -i Hardware
show arp
```

**Note**

Note the MAC addresses returned by the commands above and compare to the prior `grep` and `arp` commands and see that both hosts appear to each other as though they are on the same L2 broadcast domain. **For a little extra fun**, as you are running the pings from **host1**, on another set of windows for **s1-leaf1** & **s1-leaf2** run `clear counters` then run `watch 1`



diff show int e4 counter and see how MLAG hashing across the different pings causes the packets to choose a particular member of the port-channel in both the outbound & inbound ping flows.

## 7. Verification – on **s1-leaf1/2** and **s1-leaf3/4**

### 1. Verify the MAC addresses and the associated VTEP IP

```
s1-leaf1#show vxlan vtep
Remote VTEPS for Vxlan1:
```

VTEP	Tunnel Type(s)
10.111.253.3	unicast, flood

Total number of remote VTEPS: 1

```
s1-leaf1#show vxlan address-table
Vxlan Mac Address Table
```

VLAN	Mac Address	Type	Prt	VTEP	Moves	Last Move
112	001c.73c0.c617	DYNAMIC	Vx1	10.111.253.3	1	0:01:13 ago

Total Remote Mac Addresses for this criterion: 1

```
s1-leaf4(config)#show vxlan vtep
Remote VTEPS for Vxlan1:
```

VTEP	Tunnel Type(s)
10.111.253.1	unicast, flood

Total number of remote VTEPS: 1

```
s1-leaf4(config)#show vxlan address-table
Vxlan Mac Address Table
```

VLAN	Mac Address	Type	Prt	VTEP	Moves	Last Move
112	001c.73c0.c616	DYNAMIC	Vx1	10.111.253.1	1	0:00:33 ago

Total Remote Mac Addresses for this criterion: 1

### Note

For `show vxlan vtep` & `show vxlan address-table` to be populated, the above pings need to have been active very recently so that the MAC addresses don't age out, and you'll notice that at least 1 (but not necessarily both) of the MLAG pair switches (**s1-leaf1** or **s1-leaf2**) will have knowledge of the remote VTEP. This is because this is the direction the pings (inbound & outbound) last hashed.

## 2. Verify the MAC address and the associated interface

```
show mac address-table
```

8. Let's run some other show commands and tests to poke around VxLAN. On **s1-leaf1** and **s1-leaf4** issue the following commands:

```
show interface vxlan 1
show mac address-table
show log
```

## LAB COMPLETE!

© Copyright 2024, Arista Networks.  
Created using Sphinx (<http://sphinx-doc.org/>) 6.1.3.

[Back to top](#)